



ECOLE  
POLYTECHNIQUE  
DE BRUXELLES

MA1 - IRIF

---

HEURISTIC OPTIMIZATION

# IMPLEMENTATION EXERCISE 1

---

[https://github.com/iQuad427/PSPF\\_WT](https://github.com/iQuad427/PSPF_WT)

**Author**

Quentin ROELS

**Course**

INFO-H413

**Professor**

Thomas STUTZLE  
Christian CAMACHO  
VILLALON

**Academic Year**

2022-2023

## Abstract

This report focus on solving the Permutation Flow-shop Scheduling Problem with Weighted Tardiness (PFSP-WT) using simple Iterative Improvement algorithms and comparing their efficiency and solution quality. The second part examine the results of two Variable Neighbourhood Descent algorithms and compare them with the previous results of II algorithms. Comparisons where made using the paired Wilcoxon test to show the statistical difference between the solutions of those algorithms.

Experimental results have shown that for single neighbourhoods algorithms solution quality mainly depends on the neighbourhood, the exchange one being the best one.

For the multiple neighbourhood algorithms tested, they were unexpectedly less good than expected, seemingly due to the chosen order of neighbourhoods causing the state to rapidly find a local optima far away from the best known-solution.

## Table of contents

|          |   |           |
|----------|---|-----------|
| <b>1</b> | <b>Introduction</b>                         | <b>3</b>  |
| 1.1      | Introduction . . . . .                      | 3         |
| 1.2      | Specifications . . . . .                    | 3         |
| 1.3      | Objectives . . . . .                        | 4         |
| 1.4      | Method . . . . .                            | 4         |
| <b>2</b> | <b>Iterative Improvement</b>                | <b>5</b>  |
| 2.1      | Statistics . . . . .                        | 5         |
| 2.2      | Statistical Tests . . . . .                 | 6         |
| 2.2.1    | Initial Solution . . . . .                  | 6         |
| 2.2.2    | Pivoting Rule . . . . .                     | 7         |
| 2.2.3    | Neighbourhood . . . . .                     | 8         |
| <b>3</b> | <b>Variable Neighbourhood Descent</b>       | <b>10</b> |
| 3.1      | Comparison with II algorithms . . . . .     | 10        |
| 3.2      | Comparison between VND algorithms . . . . . | 12        |
| 3.2.1    | Initialisation . . . . .                    | 12        |
| 3.2.2    | Pivoting . . . . .                          | 12        |
| 3.2.3    | Neighbourhood . . . . .                     | 13        |
| <b>4</b> | <b>Conclusion</b>                           | <b>14</b> |
| 4.1      | Iterative Improvement . . . . .             | 14        |
| 4.2      | Variable Neighbourhood Descent . . . . .    | 14        |

## 1

## Introduction

## 1.1 Introduction

The Permutation Flow-shop Scheduling Problem (PFSP) is an active research subject with already many results. Many variants of this Operations Research classical problem exist, such as flow-time minimisation or total tardiness minimisation. One possible explanation of this popularity might be the direct application to the manufacturing industry and thus the underlying will to optimise the production lines. This problem being NP-Hard, there are no exact algorithms to solve it efficiently, allowing for continuous optimisation using new algorithms. Furthermore, due to the amount of existing results, it also makes it easy to compare the efficiency of different algorithms and the quality of the solution they provide.

## 1.2 Specifications

Focused on the Weighted Tardiness variant of the PFSP, the report investigate some of the simplest possible algorithm to optimise this kind of problem, i.e. iterative improvements algorithms. As stated earlier, the amount of previous results will allow for better comparison of the different algorithms solution. The paired Wilcoxon test has been used to show the statistical difference between the results using their deviation from the current best known solutions for given instances.

## 1.3 Objectives

This implementation exercise is divided in two parts.

- Implement simple iterative improvement algorithms to compare their solutions and execution. Those algorithms being combinations of different initialisation, pivoting and neighbourhood techniques, the goal is to see if a certain combination is better than the others.
- Implement two variable neighbourhood descent algorithms and compare them with their simple iterative improvement counterparts, to see if they lead to significantly better solutions.

## 1.4 Method

The 12 different iterative improvement and 4 variable neighbourhood descent algorithms were implemented using combinations of :

- 2 initialisation methods : random and simple RZ heuristic
- 2 pivoting rules : first improvement and best improvement
- 3 neighbourhood relations : exchange, insert and transpose

*Note : see presentation slides for more details about those*

As stated earlier, to show statistically the difference between the results of two algorithms, the paired Wilcoxon test is applied. All statistical tests are based on a null hypothesis that will either be rejected or not. For the Wilcoxon test, the null hypothesis is that "the median of the differences [between the two datasets] is zero". Thus, in this case, the test shows whether there is a significant difference between the results.

If the hypothesis is correctly rejected, it means that there is a statistical difference. This is why the test result in a p-value which indicates the probability of incorrectly rejecting the null hypothesis. This probability must be as low as possible to show that there is a difference.

Typically, a significance level  $\alpha$  is chosen (a maximum allowable value) to compare the p-value with (here  $\alpha = 0.05$ ); if the p-value is under this value, it means that this probability of falsely rejecting the null hypothesis is accepted.

In short, if a p-value is under 0.05 in this report, it means that, statistically, there is a significant difference between the values that are compared.

## Iterative Improvement

### 2.1 Statistics

Hereunder are two tables containing different statistics on the 12 algorithms that were implemented for the exercise discriminated by the size of the instance they were tested on. Table 2.1 contains the statistics for all the algorithms for instance of size 50, while Table 2.2 contains the statistics for instance of size 100.

*Note : two algorithms having the same specifications are on the same line.*

| Algorithm         | Size | Deviation | Time (sec) | Algorithm         | Size | Deviation | Time (sec) |
|-------------------|------|-----------|------------|-------------------|------|-----------|------------|
| ii-best-rnd-ex    | 50   | -29,9568  | 0,0981333  | ii-best-rnd-ex    | 100  | -15,3481  | 1,06339    |
| ii-best-rnd-ins   | 50   | -4,89609  | 0,122618   | ii-best-rnd-ins   | 100  | 0,196022  | 1,05505    |
| ii-best-rnd-tran  | 50   | 271,414   | 0,0393053  | ii-best-rnd-tran  | 100  | 87,7223   | 0,616704   |
| ii-first-rnd-ex   | 50   | -39,2156  | 0,322302   | ii-first-rnd-ex   | 100  | -23,1732  | 5,15766    |
| ii-first-rnd-ins  | 50   | -21,6357  | 0,503399   | ii-first-rnd-ins  | 100  | -9,21661  | 6,35600    |
| ii-first-rnd-tran | 50   | 264,235   | 0,0480356  | ii-first-rnd-tran | 100  | 84,6680   | 0,792992   |
| ii-best-srz-ex    | 50   | -30,8068  | 0,0906969  | ii-best-srz-ex    | 100  | -16,7051  | 0,482269   |
| ii-best-srz-ins   | 50   | -15,2437  | 0,0746060  | ii-best-srz-ins   | 100  | -13,0212  | 0,425234   |
| ii-best-srz-tran  | 50   | 85,1007   | 0,0088318  | ii-best-srz-tran  | 100  | 0,428082  | 0,0680776  |
| ii-first-srz-ex   | 50   | -40,7682  | 0,228158   | ii-first-srz-ex   | 100  | -20,8418  | 0,848928   |
| ii-first-srz-ins  | 50   | -26,7658  | 0,257870   | ii-first-srz-ins  | 100  | -14,5708  | 0,618241   |
| ii-first-srz-tran | 50   | 85,2780   | 0,0070813  | ii-first-srz-tran | 100  | 0,356621  | 0,0590493  |

Table 2.1: Algorithms statistics for 50 jobs

Table 2.2: Algorithms statistics for 100 jobs

## 2.2 Statistical Tests

This section discusses whether some algorithm components lead to better quality and/or better efficiency using the paired Wilcoxon statistical test to show if there is a statistically significant difference between the solutions.

### 2.2.1 Initial Solution

The goal of this experiment is to compare the initialisation rules, it is thus sufficient to consider one by one only the algorithms having the same specifications except for the initialisation rule.

For clarity and better understanding of the p-value result, the tables also contain the mean value of the deviation of each specification on the specific instance size.

| <b>SimpleRZ</b>   | <b>SRZ.dev</b> | <b>Random</b>     | <b>Random.dev</b> | <b>p-value</b> |
|-------------------|----------------|-------------------|-------------------|----------------|
| ii-best-srz-ex    | -30,8068       | ii-best-rnd-ex    | -29,9568          | 0,695313       |
| ii-best-srz-ins   | -15,2437       | ii-best-rnd-ins   | -4,89609          | 0,193359       |
| ii-best-srz-tran  | 85,1007        | ii-best-rnd-tran  | 271,414           | 0,00195313     |
| ii-first-srz-ex   | -40,7682       | ii-first-rnd-ex   | -39,2156          | 0,160156       |
| ii-first-srz-ins  | -26,7658       | ii-first-rnd-ins  | -21,6357          | 0,00195313     |
| ii-first-srz-tran | 85,2780        | ii-first-rnd-tran | 264,235           | 0,00195313     |

Table 2.3: Comparison of the initialisation rules for instance of size 50

| <b>SimpleRZ</b>   | <b>SRZ.dev</b> | <b>Random</b>     | <b>Random.dev</b> | <b>p-value</b> |
|-------------------|----------------|-------------------|-------------------|----------------|
| ii-best-srz-ex    | -16,7051       | ii-best-rnd-ex    | -15,3481          | 0,0644531      |
| ii-best-srz-ins   | -13,0212       | ii-best-rnd-ins   | 0,196022          | 0,00195313     |
| ii-best-srz-tran  | 0,428082       | ii-best-rnd-tran  | 87,7223           | 0,00195313     |
| ii-first-srz-ex   | -20,8418       | ii-first-rnd-ex   | -23,1732          | 0,00195313     |
| ii-first-srz-ins  | -14,5708       | ii-first-rnd-ins  | -9,21661          | 0,00195313     |
| ii-first-srz-tran | 0,356621       | ii-first-rnd-tran | 84,6680           | 0,00195313     |

Table 2.4: Comparison of the initialisation rules for instance of size 100

For both instance sizes, the p-values show that there is a significant difference in the result when using different initialisation rules, except for the exchange neighbourhood. For the transpose and insert neighbourhood, the Simple RZ initialisation tends to have better solution quality.

From the statistical result and since no difference has been shown for the exchange neighbourhood, even if the exchange neighbourhood is chosen in the next section, the Simple RZ heuristic rule seem to be the best initialisation rule.

### 2.2.2 Pivoting Rule

As for the previous section, the comparison has to be done on the same specifications except for the part of the algorithm that is under investigation. The following tables are thus separated between best and first pivoting rule.

| <b>Specs</b> | <b>Best.dev</b> | <b>Best.time</b> | <b>First.dev</b> | <b>First.time</b> | <b>dev.p-value</b> | <b>time.p-value</b> |
|--------------|-----------------|------------------|------------------|-------------------|--------------------|---------------------|
| ii-rnd-ex    | -29,9568        | 0,0981333        | -39,2156         | 0,322302          | 0,00195313         | 0,00195313          |
| ii-rnd-ins   | -4,89609        | 0,122618         | -21,6357         | 0,503399          | 0,00195313         | 0,00195313          |
| ii-rnd-tran  | 271,414         | 0,0393053        | 264,235          | 0,0480356         | 0,275391           | 0,0195313           |
| ii-srz-ex    | -30,8068        | 0,0906969        | -40,7682         | 0,228158          | 0,00195313         | 0,00195313          |
| ii-srz-ins   | -15,2437        | 0,0746060        | -26,7658         | 0,257870          | 0,00585938         | 0,00195313          |
| ii-srz-tran  | 85,1007         | 0,00883182       | 85,2780          | 0,00708134        | 0,583882           | 0,193359            |

Table 2.5: Comparison of the pivoting rules for instance of size 50

| <b>Specs</b> | <b>Best.dev</b> | <b>Best.time</b> | <b>First.dev</b> | <b>First.time</b> | <b>dev.p-value</b> | <b>time.p-value</b> |
|--------------|-----------------|------------------|------------------|-------------------|--------------------|---------------------|
| ii-rnd-ex    | -15,3481        | 1,06339          | -23,1732         | 5,15766           | 0,00195313         | 0,00195313          |
| ii-rnd-ins   | 0,196022        | 1,05505          | -9,21661         | 6,35600           | 0,00195313         | 0,00195313          |
| ii-rnd-tran  | 87,7223         | 0,616704         | 84,6680          | 0,792992          | 0,130859           | 0,00390625          |
| ii-srz-ex    | -16,7051        | 0,482269         | -20,8418         | 0,848928          | 0,00195313         | 0,00195313          |
| ii-srz-ins   | -13,0212        | 0,425234         | -14,5708         | 0,618241          | 0,0488281          | 0,00390625          |
| ii-srz-tran  | 0,428082        | 0,0680776        | 0,356621         | 0,0590493         | 0,674987           | 0,0371094           |

Table 2.6: Comparison of the pivoting rules for instance of size 100

For the deviation, the Wilcoxon test shows a statistical difference to the advantage of first improvement pivoting rule except for the algorithms using the transpose neighbourhood. This observation holds for both instance sizes.

However, for the execution time, the advantage clearly goes to the best improvement neighbourhood on both instances size. This is probably due to a faster convergence towards the optimal solution when going for the the highest improving solution.

### 2.2.3 Neighbourhood

This section compares the three neighbourhood rules, as a consequence, the comparison must be done for all pairs of neighbourhood to ensure correct observations. For better visualisation, the statistical result are divided in two; one sub-section for the deviation, and an other for the execution time.

#### Deviation

| Specs        | ins.dev  | ex.dev   | tran.dev | ins.ex.p-value | ins.tran.p-value | ex.tran.p-value |
|--------------|----------|----------|----------|----------------|------------------|-----------------|
| ii-best-rnd  | -4,89609 | -29,9568 | 271,414  | 0,00195313     | 0,00195313       | 0,00195313      |
| ii-best-srz  | -15,2437 | -30,8068 | 85,1007  | 0,00585938     | 0,00195313       | 0,00195313      |
| ii-first-rnd | -21,6357 | -39,2156 | 264,235  | 0,00195313     | 0,00195313       | 0,00195313      |
| ii-first-srz | -26,7658 | -40,7682 | 85,2780  | 0,00585938     | 0,00195313       | 0,00195313      |

Table 2.7: Comparison of the solution quality neighbourhoods rules for instance of size 50

| Specs        | ins.dev  | ex.dev   | tran.dev | ins.ex.p-value | ins.tran.p-value | ex.tran.p-value |
|--------------|----------|----------|----------|----------------|------------------|-----------------|
| ii-best-rnd  | 0,196022 | -15,3481 | 87,7223  | 0,00195313     | 0,00195313       | 0,00195313      |
| ii-best-srz  | -13,0212 | -16,7051 | 0,428082 | 0,00585938     | 0,00195313       | 0,00195313      |
| ii-first-rnd | -9,21661 | -23,1732 | 84,6680  | 0,00195313     | 0,00195313       | 0,00195313      |
| ii-first-srz | -14,5708 | -20,8418 | 0,356621 | 0,00195313     | 0,00195313       | 0,00195313      |

Table 2.8: Comparison of the solution quality of the neighbourhood rules for instance of size 100

From the statistical tests result, it immediately follows that there is a clear ranking between the different neighbourhood rules, which is the same for both instance sizes:

Exchange > Insertion > Transpose.

As a matter of fact, when comparing by pairs we see that the exchange neighbourhood is better than all the others. Then, the insertion neighbourhood is better than the transpose neighbourhood and worst than the exchange neighbourhood. Finally, the transpose neighbourhood loses against all the others



**Time**

| <b>Specs</b> | <b>ins.time</b> | <b>ex.time</b> | <b>tran.time</b> | <b>ins.ex.p-value</b> | <b>ins.tran.p-value</b> | <b>ex.tran.p-value</b> |
|--------------|-----------------|----------------|------------------|-----------------------|-------------------------|------------------------|
| ii-best-rnd  | 0,122618        | 0,0981333      | 0,0393053        | 0,00195313            | 0,00195313              | 0,00195313             |
| ii-best-srz  | 0,0746060       | 0,090697       | 0,0088318        | 0,0371094             | 0,00195313              | 0,00195313             |
| ii-first-rnd | 0,503399        | 0,322302       | 0,0480356        | 0,00195313            | 0,00195313              | 0,00195313             |
| ii-first-srz | 0,257870        | 0,228158       | 0,0070813        | 0,105469              | 0,00195313              | 0,00195313             |

Table 2.9: Comparison of the execution time of the neighbourhoods rules for instance of size 50

| <b>Specs</b> | <b>ins.time</b> | <b>ex.time</b> | <b>tran.time</b> | <b>ins.ex.p-value</b> | <b>ins.tran.p-value</b> | <b>ex.tran.p-value</b> |
|--------------|-----------------|----------------|------------------|-----------------------|-------------------------|------------------------|
| ii-best-rnd  | 1,05505         | 1,06339        | 0,616704         | 0,625000              | 0,00195313              | 0,00195313             |
| ii-best-srz  | 0,425234        | 0,482269       | 0,0680776        | 0,232422              | 0,00195313              | 0,00195313             |
| ii-first-rnd | 6,35600         | 5,15766        | 0,792992         | 0,00195313            | 0,00195313              | 0,00195313             |
| ii-first-srz | 0,618241        | 0,848928       | 0,0590493        | 0,00195313            | 0,00195313              | 0,00195313             |

Table 2.10: Comparison of the execution time of the neighbourhood rules for instance of size 100

For both instance sizes, the transpose neighbourhood has the lowest execution time of all. Then for instances of size 50, the insertion neighbourhood is slower than the exchange neighbourhood. However, for instances of size 100, the exchange neighbourhood tend to slower than the insertion neighbourhood (the difference between the two is less significant).

We can deduct that the exchange neighbourhood grows bigger than the insertion one, meaning that for bigger instances, the exchange neighbourhood will be slower, leading to this ranking :

Transpose > Insertion > Exchange.

## Variable Neighbourhood Descent

### 3.1 Comparison with II algorithms

| VND                       | Size | Deviation | Time     | ins.improve | ex.improve |
|---------------------------|------|-----------|----------|-------------|------------|
| vnd-best-rnd-tran-ex-ins  | 50   | -34,1890  | 0,277736 | 29,5733     | 5,92599    |
| vnd-best-rnd-tran-ins-ex  | 50   | -31,8161  | 0,363580 | 27,5512     | 2,74154    |
| vnd-best-srz-tran-ex-ins  | 50   | -35,6639  | 0,203165 | 22,5721     | 7,18641    |
| vnd-best-srz-tran-ins-ex  | 50   | -29,2593  | 0,197040 | 14,4490     | -2,56342   |
| vnd-first-rnd-tran-ex-ins | 50   | -37,1886  | 0,447119 | 19,6956     | -3,43519   |
| vnd-first-rnd-tran-ins-ex | 50   | -31,7998  | 0,594070 | 12,7572     | -12,4919   |
| vnd-first-srz-tran-ex-ins | 50   | -35,3642  | 0,248080 | 11,3151     | -9,21925   |
| vnd-first-srz-tran-ins-ex | 50   | -33,3451  | 0,344675 | 8,39968     | -12,6304   |
| vnd-best-rnd-tran-ex-ins  | 100  | -16,6196  | 4,34319  | 16,6164     | 1,44838    |
| vnd-best-rnd-tran-ins-ex  | 100  | -13,3810  | 4,63591  | 13,4383     | -2,33989   |
| vnd-best-srz-tran-ex-ins  | 100  | -17,4245  | 1,39865  | 4,97218     | 0,825950   |
| vnd-best-srz-tran-ins-ex  | 100  | -17,1769  | 1,43430  | 4,60941     | 0,519038   |
| vnd-first-rnd-tran-ex-ins | 100  | -21,1600  | 8,83566  | 13,1041     | -2,62463   |
| vnd-first-rnd-tran-ins-ex | 100  | -13,9024  | 10,2188  | 5,12136     | -12,0931   |
| vnd-first-srz-tran-ex-ins | 100  | -20,6418  | 1,44698  | 7,05744     | -0,250512  |
| vnd-first-srz-tran-ins-ex | 100  | -16,0205  | 1,54953  | 1,67390     | -6,09074   |

Table 3.1: Improvement of the VND algorithms over solution of previously defined II algorithms

This experiment shows that even though the VND algorithms implemented are typically better than the simple insertion neighbourhood algorithms, they seem to be less good than those with exchange neighbourhood.

It is also possible to see some other behaviour in those results :

- transpose/exchange/insert neighbourhoods tend to have better results than transpose/insert/exchange neighbourhood.
- the deviation results tend to be better when the VND is using the first improvement pivoting rule, but the improvement compared to simple neighbourhood, is better with the best improvement pivoting rule.
- VND results are closer to the simple exchange neighbourhood algorithms for bigger instances than they are for smaller instances

Those observations follow the fact that exchange neighbourhoods is typically better at finding good solutions than the insert and transpose ones. Furthermore, the quality increase caused by the first-improvement pivoting rule tend to demonstrate that slower convergence allows for better solution quality in this setting.

However, there is a lack of improvement compared to the simple exchange neighbourhood. It could mean that the transpose neighbourhood at the first place of the VND neighbourhood search tend to put the algorithm in a local optima. Which could explain the bad results of the more complex VND algorithms compared to simpler algorithms with exchange neighbourhood.

## 3.2 Comparison between VND algorithms

For the following comparisons, the same method as for comparing the iterative improvement algorithms has been used, the way the conclusions were drawn being exactly the same as before.

### 3.2.1 Initialisation

| SimpleRZ                  | SRZ.dev  | Random                    | Random.dev | p-value  |
|---------------------------|----------|---------------------------|------------|----------|
| vnd-best-srz-tran-ex-ins  | -35,6639 | vnd-best-rnd-tran-ex-ins  | -34,1890   | 0,431641 |
| vnd-best-srz-tran-ins-ex  | -29,2593 | vnd-best-rnd-tran-ins-ex  | -31,8161   | 0,160156 |
| vnd-first-srz-tran-ex-ins | -35,3642 | vnd-first-rnd-tran-ex-ins | -37,1886   | 0,375000 |
| vnd-first-srz-tran-ins-ex | -33,3451 | vnd-first-rnd-tran-ins-ex | -31,7998   | 0,492188 |

Table 3.2: Comparison of the initialisation rules for instance of size 50

| SimpleRZ                  | SRZ.dev  | Random                    | Random.dev | p-value   |
|---------------------------|----------|---------------------------|------------|-----------|
| vnd-best-srz-tran-ex-ins  | -17,4245 | vnd-best-rnd-tran-ex-ins  | -16,6196   | 0,431641  |
| vnd-best-srz-tran-ins-ex  | -17,1769 | vnd-best-rnd-tran-ins-ex  | -13,3810   | 0,0371094 |
| vnd-first-srz-tran-ex-ins | -20,6418 | vnd-first-rnd-tran-ex-ins | -21,1600   | 0,232422  |
| vnd-first-srz-tran-ins-ex | -16,0205 | vnd-first-rnd-tran-ins-ex | -13,9024   | 0,0371094 |

Table 3.3: Comparison of the initialisation rules for instance of size 100

This experiment shows that both pivoting rules give substantially the same results for both instance sizes. Except for the tran-ins-ex neighbourhood that benefits from the simple RZ heuristic for bigger instances.

### 3.2.2 Pivoting

| Best                     | Best.dev | First                     | First.dev | p-value    |
|--------------------------|----------|---------------------------|-----------|------------|
| vnd-best-rnd-tran-ex-ins | -34,1890 | vnd-first-rnd-tran-ex-ins | -37,1886  | 0,0273438  |
| vnd-best-rnd-tran-ins-ex | -31,8161 | vnd-first-rnd-tran-ins-ex | -31,7998  | 0,845703   |
| vnd-best-srz-tran-ex-ins | -35,6639 | vnd-first-srz-tran-ex-ins | -35,3642  | 0,921875   |
| vnd-best-srz-tran-ins-ex | -29,2593 | vnd-first-srz-tran-ins-ex | -33,3451  | 0,00390625 |

Table 3.4: Comparison of the pivoting rules for instance of size 50

| <b>Best</b>              | <b>Best.dev</b> | <b>First</b>              | <b>First.dev</b> | <b>p-value</b> |
|--------------------------|-----------------|---------------------------|------------------|----------------|
| vnd-best-rnd-tran-ex-ins | -16,6196        | vnd-first-rnd-tran-ex-ins | -21,1600         | 0,00195313     |
| vnd-best-rnd-tran-ins-ex | -13,3810        | vnd-first-rnd-tran-ins-ex | -13,9024         | 0,275391       |
| vnd-best-srz-tran-ex-ins | -17,4245        | vnd-first-srz-tran-ex-ins | -20,6418         | 0,00585938     |
| vnd-best-srz-tran-ins-ex | -17,1769        | vnd-first-srz-tran-ins-ex | -16,0205         | 0,556641       |

Table 3.5: Comparison of the pivoting rules for instance of size 100

The only algorithms and instance size for which there is statistical difference are the one for which the first improvement pivoting rule is better. This experiment shows that the first improvement pivoting rule is statistically better than the best improvement rule.

### 3.2.3 Neighbourhood

| <b>ins.ex</b>             | <b>ins.ex.dev</b> | <b>ex.ins</b>             | <b>ex.ins.dev</b> | <b>p-value</b> |
|---------------------------|-------------------|---------------------------|-------------------|----------------|
| vnd-best-rnd-tran-ins-ex  | -31,8161          | vnd-best-rnd-tran-ex-ins  | -34,1890          | 0,193359       |
| vnd-best-srz-tran-ins-ex  | -29,2593          | vnd-best-srz-tran-ex-ins  | -35,6639          | 0,00195313     |
| vnd-first-rnd-tran-ins-ex | -31,7998          | vnd-first-rnd-tran-ex-ins | -37,1886          | 0,00195313     |
| vnd-first-srz-tran-ins-ex | -33,3451          | vnd-first-srz-tran-ex-ins | -35,3642          | 0,275391       |

Table 3.6: Comparison of the neighbourhood rules for instance of size 50

| <b>ins.ex</b>             | <b>ins.ex.dev</b> | <b>ex.ins</b>             | <b>ex.ins.dev</b> | <b>p-value</b> |
|---------------------------|-------------------|---------------------------|-------------------|----------------|
| vnd-best-rnd-tran-ins-ex  | -13,3810          | vnd-best-rnd-tran-ex-ins  | -16,6196          | 0,00195313     |
| vnd-best-srz-tran-ins-ex  | -17,1769          | vnd-best-srz-tran-ex-ins  | -17,4245          | 1,00000        |
| vnd-first-rnd-tran-ins-ex | -13,9024          | vnd-first-rnd-tran-ex-ins | -21,1600          | 0,00195313     |
| vnd-first-srz-tran-ins-ex | -16,0205          | vnd-first-srz-tran-ex-ins | -20,6418          | 0,00390625     |

Table 3.7: Comparison of the neighbourhood rules for instance of size 100

The only algorithms and instance size for which there is statistical difference are the one for which the transpose-exchange-insertion neighbourhood is better. This experiment shows that the "transpose, exchange, insertion" neighbourhood is statistically better than the "transpose, insertion, exchange" one.

## 4

## Conclusion

## 4.1 Iterative Improvement

When comparing the different combinations of operators for the iterative improvement algorithm, the results tend to show that the combination of either of the two initialisation methods, followed by a first-improvement pivoting rules using the exchange neighbourhood will eventually lead to the best results. This conclusion can be corroborated using the Tables 2.1 and 2.2 where the best deviations measured were for the ii-first-srz-ex and ii-first-rnd-ex algorithms.

## 4.2 Variable Neighbourhood Descent

Unexpectedly, the results on the implementations of the VND algorithms did not surpass the results quality of the best combination for the Iterative Improvement algorithms. As stated earlier, this result could be explained either by the order of the neighbourhoods (starting with transpose neighbourhood) or by a flaw in the implementation of the algorithm. Seeing the result and the complementary tests made on other order of neighbourhoods, it seems like starting with the transpose neighbourhood tend to lock the execution on a local optimum state, leading to worse solution quality than the best combination of the previous section algorithms.

When comparing the VND algorithms between each other, conclusions are quite direct. While the initial solution construction does not affect the results significantly, a clear advantage goes to the first-improvement method compared to the best-improvement pivoting rule for both instance sizes. The neighbourhood order see the advantage going to the transpose/exchange/insert order compared to the transpose/insert/exchange one. This advantage grows even more with the increasing size of the instances.